

Grand Theater	
Showings	
Star Wars	7pm
Star Wars	9pm
Spectre	6:30pm

$\text{List}[\text{Event}]$   
 $\text{latency}(\text{page load}) \leq 20\text{ms}$

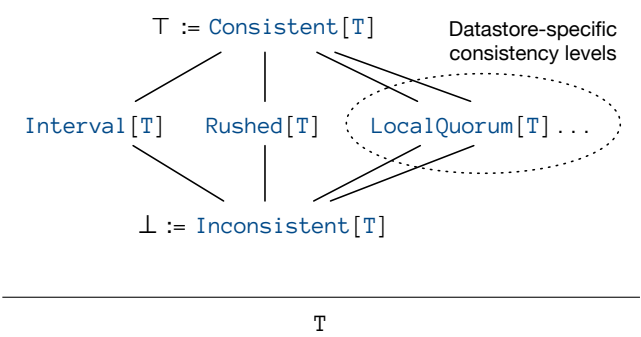
STAR WARS

7pm Remaining:

5

Purchase

Counter  
invariant:  $\text{tickets} \geq 0$     error tolerance:  $\text{tickets} \pm 5\%$



×

Counter

×

$\text{SetTSet}[T]$  with  $\text{Consistency}(\text{Strong})$   $\text{SetSet}[T]$  with  $\text{LatencyB}$

$\text{LatencyBound}(x)$   
 $\text{ErrorTolerance}(x\%)\text{sizeSet}$

$\text{Setsizecontains}$   
 $\text{ErrorTolerance}(5\%)\text{List}$

$T\text{Inconsistent}[T]\text{Consistent}[T]T \prec$

$$\frac{\tau \tau'}{\tau'[T] \prec \tau[T]}$$

BoundedCounter

$\text{Consistent}[T]T$   
 $\text{Inconsistent}[T]\text{Inconsistent}[T]\text{Consistent}(x)\text{Consistent}[T]$

$\text{Set}[T]$

$$\frac{\Gamma \vdash e_1 : \tau[T]T \prec \tau[T]}{\Gamma \vdash \text{Consistent}(e_1) : T}$$

Inconsistent[T]

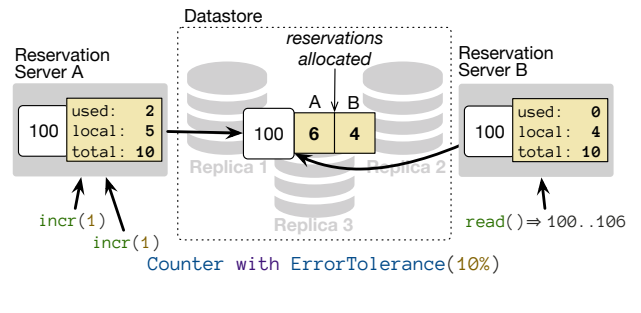
Rushed[T] LatencyBound(x)Rushed[T] LatencyBoundRushed[T]  
Rushed[T]  
Rushed[T] Consistent[T] LocalQuorum[T] Inconsistent[T]

```
set.contains() match {  
  case Consistent(x) => print(x)  
  case LocalQuorum(x) => print(x + ", locally")  
  case Inconsistent(_) => print("unknown")  
}  
  
Inconsistent(_)
```

```
Interval[T] ErrorTolerance(x%) Interval[T]  
Interval[T] Set size add size size [95, 105] [100, 107] add  
  
Interval[Int]
```

$$W + R > N(N + 1) / 2 \text{QUORUM} N \text{ALL ONE LOCAL\_QUORUM}$$

QUORUMQUORUM



```
Set add remove size  
  
Interval(min = v - removePool.delta()  
          max = v + addPool.delta())  
  
v delta  
  
delta(): pool.total - (pool.local - pool.used)
```

LatencyBound

Counter  
Set add remove contain size size  
BoundedCounter  
UUIDPoolBoundedCounter  
List

tc netemnetem

Counter  
ALL  
[100, 106]

incr(1)read

```
trait LatencyBound {
  // execute readOp with strongest consistency possible
  // within the latency bound
  def rush[T](bound: Duration,
    readOp: ConsistencyLevel => T): Rushed[T]
}



/* Generic reservaton pool, conceptually one per
 * ADT instance. `max` recomputed as needed
 * (e.g. for percent error) */
abstract class ReservationPool(max: () => Int) {
  def take(n: Int): Boolean // try to take tokens
  def sync(): Unit // sync to regain used tokens
  def delta(): Int // # possible ops outstanding
}

/* Counter with ErrorBound (simplified) */
class Counter(key: UUID) with ErrorBound {
  def error: Float // error bound
  def computeMax(): Int = (cass.read(key) * error).toInt

  val incrPool = ReservationPool(computeMax)
  val decrPool = ReservationPool(computeMax)

  def value(): Interval[Int] = {
    val v = cass.read(key)
    Interval(v - decrPool.delta,
      v + incrPool.delta)
  }

  def incr(n: Int): Unit = {
    waitFor(incrPool.take(n)) {
      cass.incr(key, n)
    }
  }
}
```






QUORUM  
weakALLQUORUM

ALL

us-eastus-centraleurope-westasia-east

browse  
viewEvent  
purchase  
addEvent

ListUUIDPoolBoundedCounter  
BoundedCounter


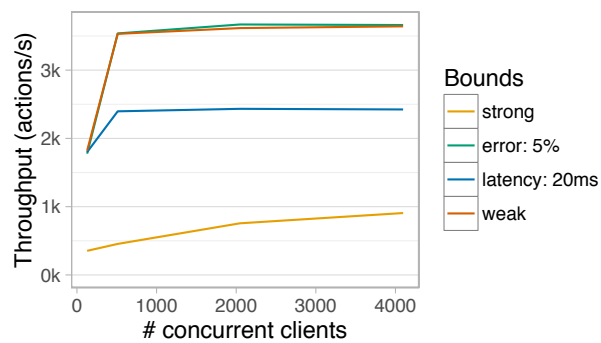


```
User(id: UserID,
  name: String,
  followers: Set[UserID] with LatencyBound(20 ms),
  timeline: List[TweetID] with LatencyBound(20 ms)
)

Tweet(id: TweetID,
  user: UserID,
  text: String,
  posted: DateTime,
  retweets: Set[UserID] with Size(ErrorTolerance(5%))
)
```

size

×Pool



×

tweet

Interval

viewEventbrowsepurchaseaddEvent×purchaseBoundedCounter××  
viewEventbrowse

PoolCounter

SetSetListListSet

	<a href="http://hyperdex.org/">http://hyperdex.org/</a>
	<a href="https://www.usenix.org/conference/osdi12/technical-sessions/presentation/li">https://www.usenix.org/conference/osdi12/technical-sessions/presentation/li</a>
Rushed[T]	<a href="http://glinden.blogspot.com/2006/12/slides-from-my-talk-at-stanford.html">http://glinden.blogspot.com/2006/12/slides-from-my-talk-at-stanford.html</a>
	<a href="https://www.usenix.org/conference/nsdi14/technical-sessions/presentation/liu_jed">https://www.usenix.org/conference/nsdi14/technical-sessions/presentation/liu_jed</a>
	<a href="http://www.wired.com/2015/11/how-instagram-solved-its-justin-bieber-problem/">http://www.wired.com/2015/11/how-instagram-solved-its-justin-bieber-problem/</a>
	<a href="http://www.buzzfeed.com/daozers/what-its-like-to-work-on-buzzfeeds-tech-team-during-record-t">http://www.buzzfeed.com/daozers/what-its-like-to-work-on-buzzfeeds-tech-team-during-record-t</a>
	<a href="http://outworkers.github.io/phantom/">http://outworkers.github.io/phantom/</a>
<a href="http://fusionticket.org">http://fusionticket.org</a>	
<a href="https://aws.amazon.com/ec2/">https://aws.amazon.com/ec2/</a>	
<a href="http://cassandra.apache.org/">http://cassandra.apache.org/</a>	
<a href="http://www.reuters.com/article/2014/03/03/us-oscars-selfie-idUSBREA220C320140303">http://www.reuters.com/article/2014/03/03/us-oscars-selfie-idUSBREA220C320140303</a>	
<a href="http://docs.basho.com/riak/latest/">http://docs.basho.com/riak/latest/</a>	
<>	<a href="http://redis.io/">http://redis.io/</a>
	<a href="http://redis.io/topics/twitter-clone">http://redis.io/topics/twitter-clone</a>
<a href="http://www.forbes.com/sites/hayleycuccinello/2015/10/20/star-wars-presales-crash-ticketing-sites-sets-record-for-fandango/">http://www.forbes.com/sites/hayleycuccinello/2015/10/20/star-wars-presales-crash-ticketing-sites-sets-record-for-fandango/</a>	
<a href="http://docs.datastax.com/en/cassandra/3.x/cassandra/dml/dmlAboutDataConsistency.html">http://docs.datastax.com/en/cassandra/3.x/cassandra/dml/dmlAboutDataConsistency.html</a>	<a href="http://www.linuxfoundation.org/collaborate/workgroups/networking/netem">http://www.linuxfoundation.org/collaborate/workgroups/networking/netem</a>
	<a href="https://twitter.github.io/finagle/">https://twitter.github.io/finagle/</a>
<a href="https://www.docker.com/">https://www.docker.com/</a>	<a href="https://www.usenix.org/conference/osdi14/technical-sessions/presentation/xie">https://www.usenix.org/conference/osdi14/technical-sessions/presentation/xie</a>
<a href="http://radar.oreilly.com/2009/06/bing-and-google-agree-slow-pag.html">http://radar.oreilly.com/2009/06/bing-and-google-agree-slow-pag.html</a>	
<a href="https://cloud.google.com/compute/">https://cloud.google.com/compute/</a>	